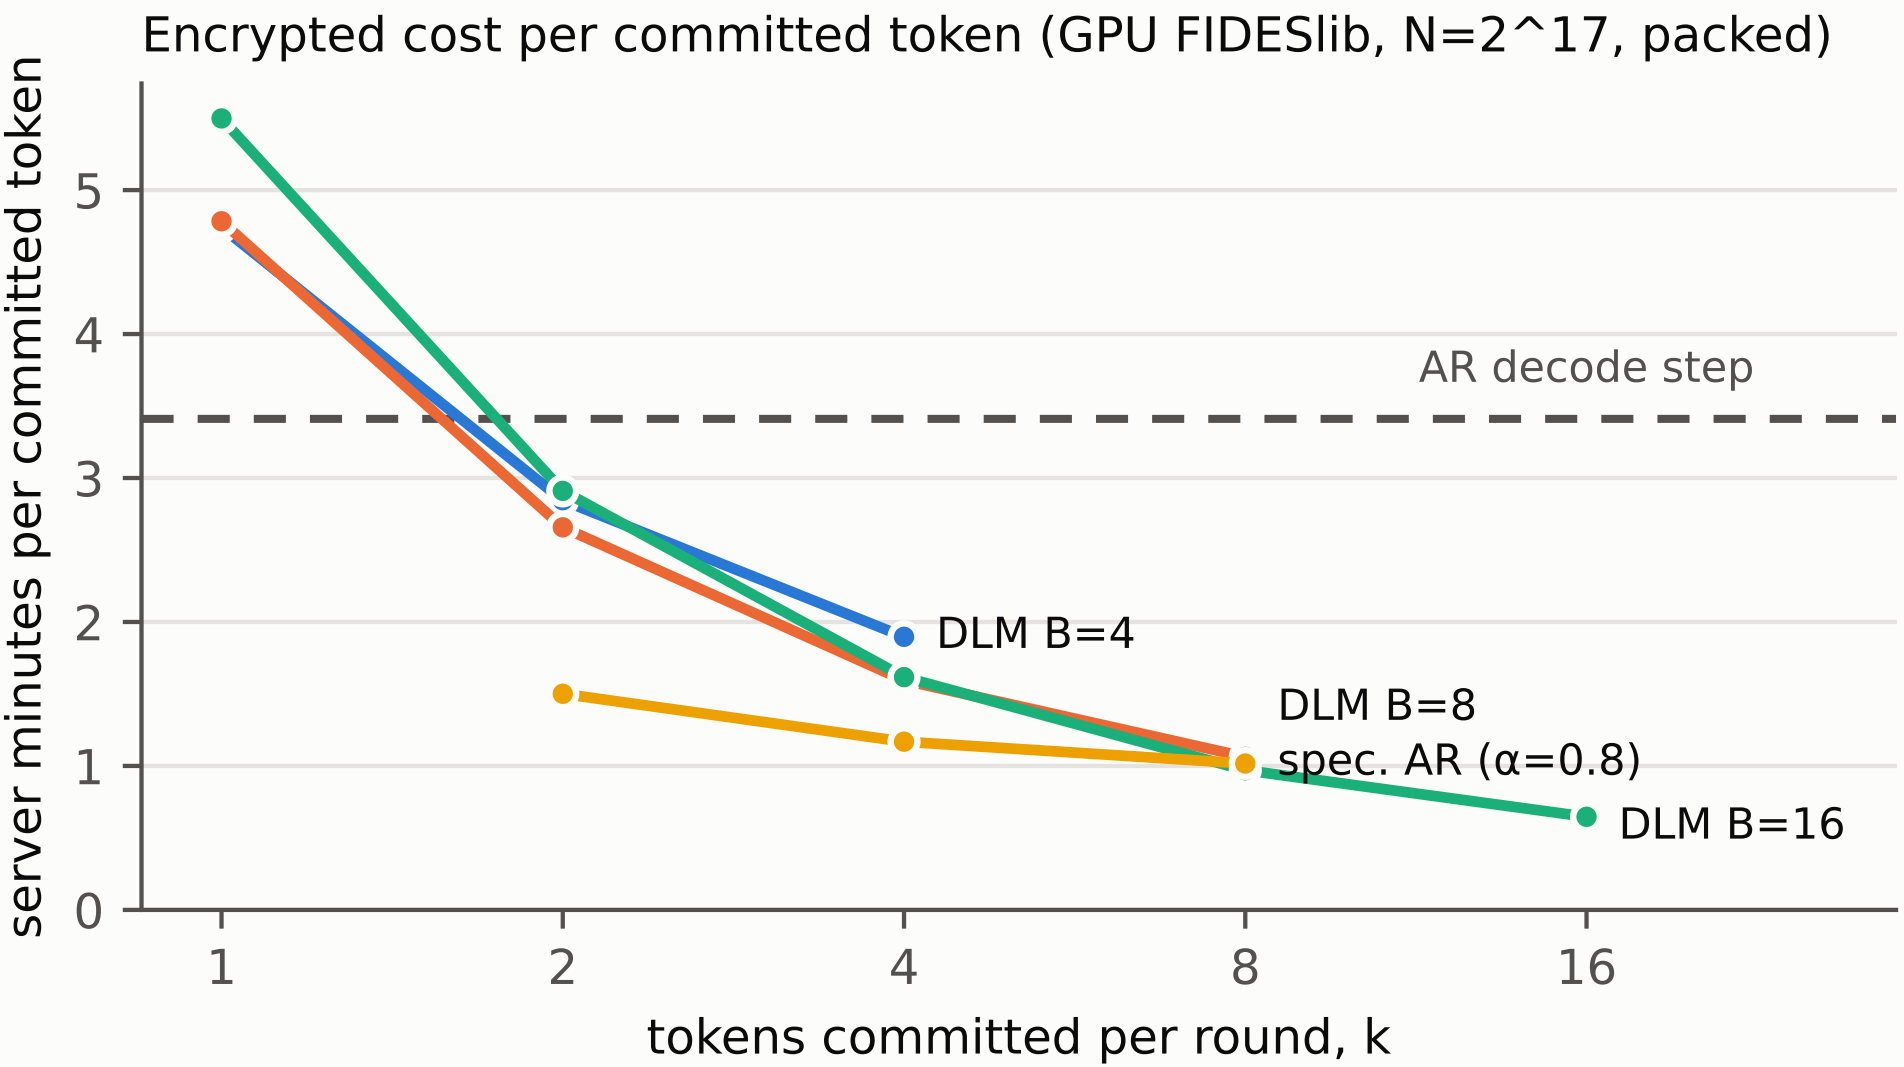

Model from measured op costs; L=256, 12 layers, d padded to 1024, GELU deg 31;
matvec measured (compute+encode). Spec. AR: k = drafts verified per pass.
GPU: bootstrap (4 GPUs) and matvec (1 GPU) measured, other ops scaled; network excluded.